

# Natural Language Processing

Lecture 11:  
Machine Learning: Feed-forward Neural Networks


10/23/2018 & 10/25/2018

COMS W4705  
Daniel Bauer


# Perceptron Expressiveness

- Simple perceptron learning algorithm, starts with an arbitrary hyperplane and adjusts it using the training data.
  - Step function is not differentiable, so no closed-form solution.
- Perceptron produces a linear separator.
  - Can only learn linearly separable patterns.
- Can represent boolean functions like **and**, **or**, **not** but not the **xor** function.


# The problem with *xor*



(a)  $x_1$  **and**  $x_2$




(b)  $x_1$  **or**  $x_2$



(c)  $x_1$  **xor**  $x_2$



# Multi-Layer Neural Networks



- Basic idea: represent any (non-linear) function as a composition of soft-threshold functions. This is a form of non-linear regression.
- Lippmann 1987: Two hidden layers suffice to represent any arbitrary region (provided enough neurons), even discontinuous functions!

# Activation Functions

- One problem with perceptrons is that the **threshold function (step function)** is undifferentiable.
- It is therefore unsuitable for gradient descent.
- One alternative is the **sigmoid (logistic) function**:




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g(z) = 0 \text{ if } z \rightarrow -\infty$$


$$g(z) = 1 \text{ if } z \rightarrow \infty$$

# Activation Functions

- Two other popular activation functions:






$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$




$$\text{relu}(z) = \max(z, 0)$$

# How Do Neural Networks Represent Functions?



# Output Representation


- Many NLP Problems are multi-class classification problems.
- Each output neuron represents one class. Predict the class with the highest activation.





# Softmax

- We often want the activation at the output layer to represent probabilities.
- Normalize activation of each output unit by the sum of all output activations (as in log-linear models).




$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}$$

The network computes a probability

$$P(c_i | \mathbf{x}; \mathbf{w})$$

# Softmax

- We often want the activation at the output layer to represent probabilities.
- Normalize activation of each output unit by the sum of all output activations (as in log-linear models).



$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}$$

The network computes a probability

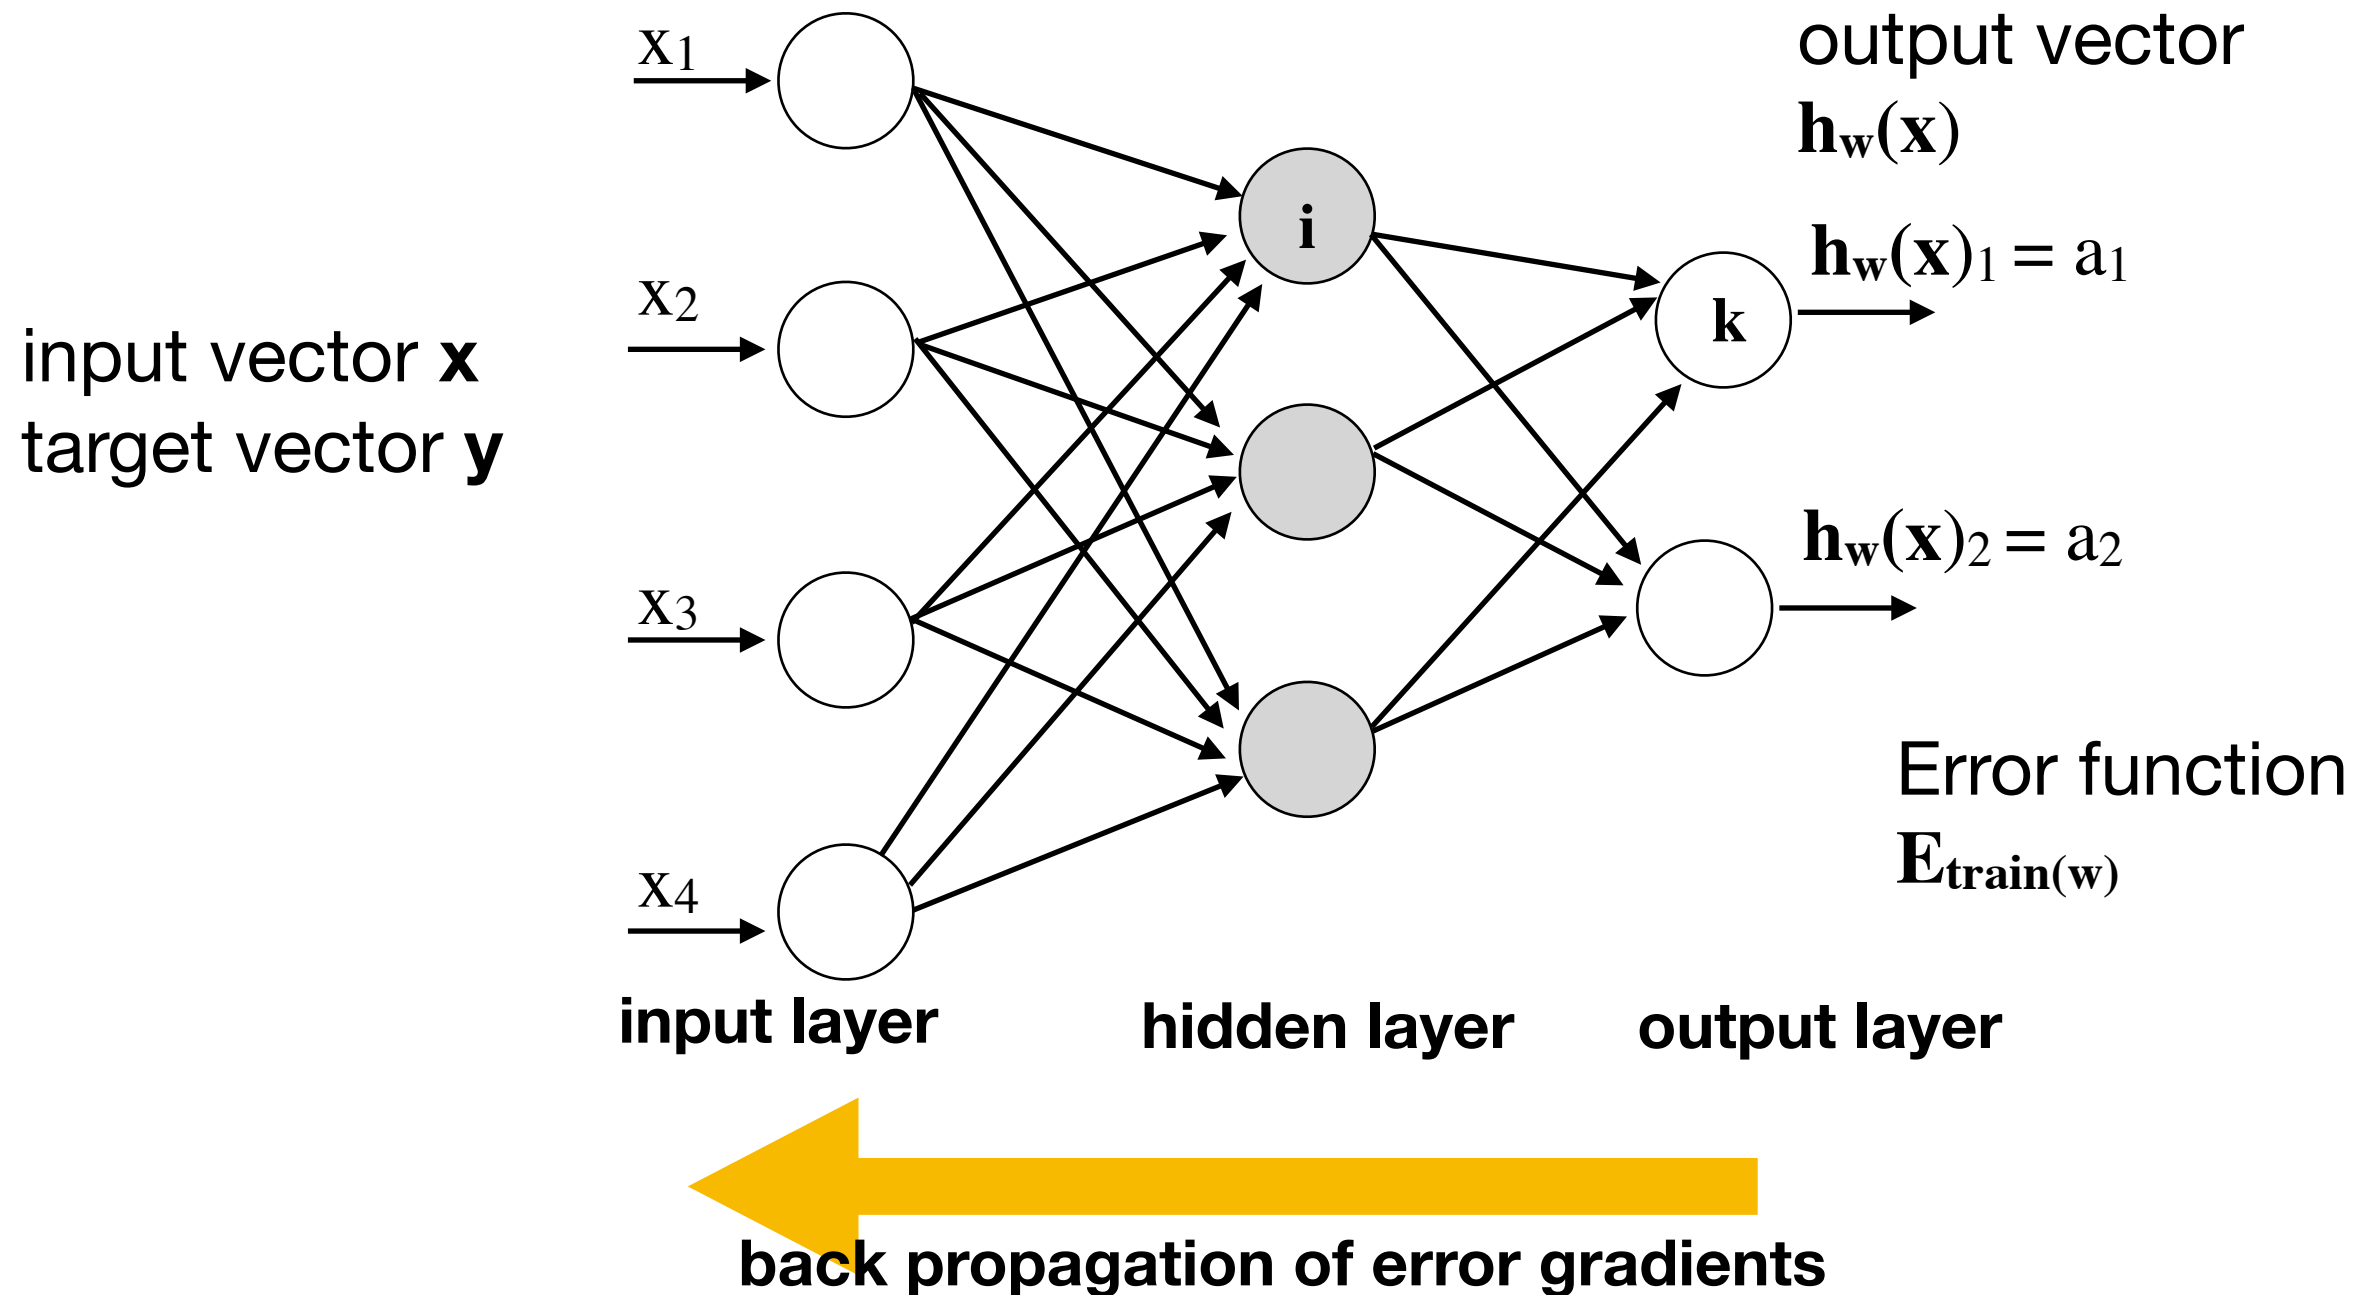
$$P(c_i | \mathbf{x}; \mathbf{w})$$

# Learning in Multi-Layer Neural Networks

- Network structure is fixed, but we want to train the weights. Assume **feed-forward** neural networks: no connections that are loops.
- **Backpropagation Algorithm:**
  - Given current weights, get network output and compute loss function (assume multiple outputs / a vector of outputs).
  - Can use gradient descent to update weights and minimize loss.
  - Problem: We only know how to do this for the last layer!
  - Idea: Propagate error backwards through the network.

# Backpropagation

feed-forward computation of network outputs



# Negative Log-Likelihood

(also known as cross-entropy)

- Assume target output is a one-hot vector and  $c(y)$  is the target class for target  $\mathbf{y}$ .
- Compute the negative log-likelihood for a single example

$$Loss(\mathbf{y}, h_{\mathbf{w}}(x)) = -\log P(c(\mathbf{y}) | \mathbf{x}; \mathbf{w})$$

- Empirical error for the entire training data:

$$E_{train}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N -\log P(c(\mathbf{y}^{(i)}) | \mathbf{x}^i; \mathbf{w})$$

# Stochastic Gradient Descent

- Goal: Learn parameters that minimize the empirical error.

Randomly initialize  $w$

for a set number of iterations  $T$ :

shuffle training data  $\mathcal{D} = (x^{(j)}, y^{(j)})|_{j=1}^n$

for  $j = 1 \dots N$ :

for each  $w_i$  (all weights in the network):

$$w_i \leftarrow w_i - \eta \frac{\partial}{\partial w_i} \text{Loss}(y^{(j)}, h_{\mathbf{w}}(x^{(j)}))$$

- $\eta$  is the learning rate.
- It often makes sense to compute the gradient over batches of examples, instead of just one ("mini-batch").

# Backpropagation

- Simplified multi-layer case (a single unit per layer):



- Stochastic Gradient Descent should perform the following update:

$$w_2 \leftarrow w_2 - \eta \frac{\partial \text{Loss}(y, f(g(x)))}{\partial w_2}$$

$$w_1 \leftarrow w_1 - \eta \frac{\partial \text{Loss}(y, f(g(x)))}{\partial w_1}$$

- Problem: How do we compute the gradient for parameters  $w_1$  and  $w_2$ ?

# Chain Rule of Calculus

- To compute gradients for hidden units, we need to use apply the chain rule of calculus:

The derivative of  $f(g(x))$  is

$$\frac{df(g(x))}{dx} = \frac{df(g(x))}{dg(x)} \cdot \frac{dg(x)}{dx}$$



# Backpropagation




$$\frac{\partial Loss}{\partial w_2} = \left( \frac{\partial Loss}{\partial f(g(x))} \right) \left( \frac{\partial f(g(x))}{\partial w_2} \right)$$

$$\frac{\partial Loss}{\partial w_1} = \left( \frac{\partial Loss}{\partial g(x)} \right) \left( \frac{\partial g(x)}{\partial w_1} \right)$$


$$= \left( \frac{\partial Loss}{\partial f(g(x))} \right) \left( \frac{\partial f(g(x))}{\partial g(x)} \right) \left( \frac{\partial g(x)}{\partial w_1} \right)$$

# Backpropagation

forward



backward




Assume we know  $\frac{\partial Loss}{\partial f(x)}$

We want to compute  $\frac{\partial Loss}{\partial x}$  to propagate it back.


and  $\frac{\partial Loss}{\partial w}$  (for the weight update)

# Backpropagation

forward



backward




$$\frac{\partial \text{Loss}}{\partial x} = \left( \frac{\partial \text{Loss}}{\partial f(x)} \right) \left( \frac{\partial f(x)}{\partial x} \right)$$

$$\frac{\partial \text{Loss}}{\partial w} = \left( \frac{\partial \text{Loss}}{\partial f(x)} \right) \left( \frac{\partial f(x)}{\partial w} \right)$$

these depend on the function  $f$ .

# Backpropagation with Multiple Neurons

- Let  $\Delta_j = \frac{\partial Loss}{\partial j}$  be the derivative of the loss w.r.t to the output of unit j.




$$\begin{aligned}\Delta_i &= \frac{\partial Loss}{\partial i} = \sum_j \left( \frac{\partial Loss}{\partial j} \right) \left( \frac{\partial j}{\partial i} \right) \\ &= \sum_j \Delta_j \left( \frac{\partial j}{\partial i} \right)\end{aligned}$$

- The output of j is computed during the forward pass.

# Backpropagation with Multiple Neurons

- Once the  $\Delta_j$  have been computed, we can compute the gradients w.r.t to the weights.



$$\frac{\partial Loss}{\partial w_{ij}} = \Delta_j \frac{\partial j}{\partial w_{ij}}$$

# Some Neural Network Tricks

- When implementing the model, try to fit to 100% accuracy on 1 or 2 data points.
- Decrease learning rate after each epoch or when loss stops decreasing.
- Find good initial learning rate before optimizing other hyperparameters.
- Try other optimizers:
  - SGD with momentum, rmsprop, adagrad, adadelata, adam

# Acknowledgments

- Some slides by Chris Kedzie